

End-to-End Flow

1. User Subscription

- User clicks “Subscribe” button in the frontend.
- Backend verifies **user permissions/rights**.
- If allowed, a **subscription record** is saved in the database.
- A **scheduled job** is also created with the necessary metadata (user ID, report type, schedule info, etc.).

2. Job Scheduling

- A **background service** periodically scans for scheduled jobs whose **createdAt** or **nextRun** timestamp is due.
- For each eligible job, it **pushes a message into a job queue** (e.g., SQS, RabbitMQ) describing the work to be done (documents to aggregate, processing details).

3. Worker Processing

- **Workers** are subscribed to the queue.
- Each worker picks up a message and performs **heavy processing**:
 - Fetching documents/images.
 - Aggregating data.
 - Performing image or PDF processing.
- Once processing completes, the **PDF report is stored in S3** (or another object store).

4. Notification & Email

- When the PDF is successfully stored, a **new message/event** is triggered for the notification process.
- A **Notification Worker** picks this up, generates a **safe link** to the PDF, and invokes the **Email Service**.
- The **user receives the report email**.

Worker Processing with Three Workers

Fetch Worker

- Picks up jobs from the queue that indicate which documents/images need to be processed.
- Fetches the data from the database and S3.
- Pushes the fetched raw data into the next queue for aggregation.

Aggregation Worker

- Picks up fetched data from the fetch queue.
- Aggregates the tabular and textual data into a coherent dataset.
- Prepares a structured payload for PDF/image generation.
- Pushes the aggregated data to the next queue for report generation.

Report/Processing Worker

- Picks up aggregated data from the aggregation queue.
- Performs **image processing**, merges tables/text, and generates the **final PDF report**.
- Stores the PDF in S3.
- Triggers the **notification queue** for safe link generation and email delivery.

Advantages of This Approach

- Each worker can **scale independently** based on workload.
- Failures in one step **don't block others**, allowing retry strategies per worker.
- Easier to **monitor and optimize** specific parts of the pipeline (fetching, aggregation, processing).

Fetch Worker → Aggregation Queue → Aggregation Worker → Processing Queue → Report Worker